

(*

Interlude: abstracting from the particular language

- to a theory of all abstract languages, and
- (all) possible interpreters (evaluators)

*)

(*

Signatures and Signed/Ranked Algebras *)

(*

Let us present a general framework for discussing Abstract Syntax

For simplicity, we confine our interest to a homogeneous algebraic framework: all expressions are of one *sort*. (The generalisation is called multi-sorted algebras.)

*)

(*

We first fix the symbols which any candidate (abstract) language will use:

Definition: Signature Σ - a set of symbols,
and for each symbol, its arity ≥ 0

*)

(* **Example:** Boolean expressions

\mathbf{T} , \mathbf{F} arity 0

\neg arity 1

\wedge , \vee arity 2

*)

(* **Example:** Integer expressions

(a denumerable set of) **numerals** arity 0

$+$, $*$ arity 2

unary minus $-$ arity 1

*)

(*

We now discuss the meaning of expressions in any language whose symbols have been specified using a signature Σ .

Semantics

Definition (Σ -algebra)

A — a carrier set, and

for each 0-ary symbol in Σ , associate some element of A

for each k -ary symbol in Σ , associate a total function in $A^k \rightarrow A$

Notes:

- Not all elements of A require having a 0-ary symbol in Σ associated with them. That is, the carrier set may contain values (perhaps infinitely many) for which there is no syntax. Irrational real numbers are a good examples — there may be

no way of writing a symbol for each irrational number. In fact, we may not even be able to write an expression for each irrational number.

- Different 0 -ary symbols in Σ do not necessarily have to be associated with different elements — two different symbols can be mapped to the same value. Similarly for k -ary symbols, two different symbols can be mapped to the same function.
- Two Σ -algebras can have the same carrier set A , but differ on the interpretation of the symbols in the signature Σ . These would be considered different Σ -algebras.

*)

(*)

Note: Symbols do not inherently carry meaning. We can assign whatever meanings we choose to but once we assign meaning by choosing a particular Σ -algebra, we have fixed the meanings of the symbols.

In particular, we can give

Standard examples of Σ -algebras, where the symbols are given meanings that we all expect and agree upon, e.g., binary symbol $+$ means addition on naturals, the 0 -ary symbol numeral 0 means the natural number zero.

But we can also give

non-standard examples of Σ -algebras, where the meanings of symbols differ from what we normally expect, e.g., 0 -ary symbol numeral 0 can be given the meaning “forty-two”.

*)

(*)

Notation:

If $\mathcal{A} = \langle A, \dots \rangle$ is a Σ -algebra, and $a \in A$ is associated with 0 -ary symbol c in Σ , we write $c_{\mathcal{A}}$ to denote the element $a \in A$, highlighting that this is the meaning associated with symbol c in algebra \mathcal{A} .

Similarly, if f is a k -ary symbol in Σ , and in Σ -algebra $\mathcal{A} = \langle A, \dots \rangle$, f is associated with a total function $g : A^k \rightarrow A$, we write $f_{\mathcal{A}}$ to denote this function g .

*)

(*) **Abstract Syntax, abstractly and formally**

Definition: (Carrier set $Tree_{\Sigma}$)

Suppose Σ is a given signature.

The set $Tree_{\Sigma}$ is inductively defined as follows:

Base cases: for each 0 -ary symbol c in Σ , a labelled node $\bullet c$ is in $Tree_{\Sigma}$

Induction cases: for each $k > 0$,

for each k -ary symbol f in Σ ,

for any trees t_1, \dots, t_k (already) in $Tree_{\Sigma}$,

the tree consisting of

a labelled node $\bullet f$ at the root

with the k trees t_1, \dots, t_k ordered as subtrees below the root node

is in $Tree_{\Sigma}$.

Notes: The roots of t_1, \dots, t_k are the children, ordered $1 \dots k$, of the new root node $\bullet f$. The t_i need not be distinct.

The Abstract Syntax Tree Algebra: *Syntax as Semantics*

Definition: \mathbf{Tree}_Σ is the Σ -algebra with carrier set \mathbf{Tree}_Σ , where

- every 0-ary symbol c in Σ is interpreted as the labelled node $\bullet c$, which is in \mathbf{Tree}_Σ , and
- every k -ary symbol f in Σ is interpreted as the *tree-forming function* that takes k trees t_1, \dots, t_k and creates the *tree*



in \mathbf{Tree}_Σ by sticking t_1, \dots, t_k as the *subtrees* below a new root node labelled $\bullet f$.

*)

(*)

Structure-preserving functions between Σ -algebras

Σ -homomorphisms are *structure-preserving functions* between Σ -algebras which respect the *symbol association* (imposed with respect to the symbols in the *signature* Σ) by the semantics of the source and target Σ -algebras.

Definition (Σ -homomorphisms) Suppose Σ is a given *signature*.

Suppose $\mathcal{A} = \langle A, \dots \rangle$ and $\mathcal{B} = \langle B, \dots \rangle$ are both Σ -algebras (for the same *signature* Σ).

A *total function* $h : A \rightarrow B$ between the *carrier sets* of these algebras is called a Σ -*homomorphism* if

- for all 0-ary symbols c in Σ : $h(c_{\mathcal{A}}) = c_{\mathcal{B}}$

- for all k -ary ($k > 0$) symbols f in Σ ,

for all $a_1, \dots, a_k \in A$,

$$h(f_{\mathcal{A}}(a_1, \dots, a_k)) = f_{\mathcal{B}}(h(a_1), \dots, h(a_k))$$

*)

(*)

Initiality Theorem

For any Σ -algebra, $\mathcal{B} = \langle B, \dots \rangle$, there is a unique Σ -homomorphism

$$i_{\mathcal{B}} : \mathbf{Tree}_\Sigma \rightarrow B$$

from the Σ -algebra \mathbf{Tree}_Σ to the Σ -algebra \mathcal{B} .

Proof — by construction, define $i_{\mathcal{B}}$ to be a Σ -homomorphism.

Define $i_{\mathcal{B}} : \mathbf{Tree}_\Sigma \rightarrow B$ as follows:

- for each 0-ary symbol c in Σ : $i_{\mathcal{B}}(\bullet c) = c_{\mathcal{B}}$

- for each k -ary ($k > 0$) symbol f in Σ ,

$$\begin{array}{c} i_{\mathcal{B}}(\bullet f) = f_{\mathcal{B}}(i_{\mathcal{B}}(t_1), \dots, i_{\mathcal{B}}(t_k)) \\ / \quad \backslash \\ t_1 \dots t_k \end{array}$$

Uniqueness of $i_{\mathcal{B}}$

Suppose $j : \text{Tree}_{\Sigma} \rightarrow \mathcal{B}$ is also a Σ -homomorphism from Tree_{Σ} to \mathcal{B} .

Proof by induction on $(\text{ht } t)$ that for all t in Tree_{Σ} , $i_{\mathcal{B}}(t) = j(t)$

Base cases $(\text{ht } t) = 0$.

t must be of the form $\bullet c$ for some 0-ary symbol in Σ .

But for each 0-ary symbol c in Σ : $i_{\mathcal{B}}(\bullet c) = c_{\mathcal{B}}$ // defn of $i_{\mathcal{B}}$
 $= j(\bullet c)$ // j is a Σ -homomorphism
// from Tree_{Σ} to \mathcal{B}

Induction Hypothesis:

Assume that for all t' in Tree_{Σ} such that $(\text{ht } t') \leq n$,

$$i_{\mathcal{B}}(t') = j(t')$$

Induction Step: Consider t in Tree_{Σ} such that $(\text{ht } t) = n+1$

Therefore t is of the form

$$\begin{array}{c} \bullet f \\ / \quad \backslash \\ t_1 \dots t_k \end{array}$$

with $(\text{ht } t_i) \leq n$ ($1 \leq i \leq k$)

Now, by definition of $i_{\mathcal{B}}$,

for each k -ary symbol ($k > 0$) f in Σ ,

$$\begin{array}{c} i_{\mathcal{B}}(\bullet f) \\ / \quad \backslash \\ t_1 \dots t_k \end{array}$$

$$= f_{\mathcal{B}}(i_{\mathcal{B}}(t_1), \dots, i_{\mathcal{B}}(t_k)) \text{ // definition of } i_{\mathcal{B}}$$

$$= f_{\mathcal{B}}(j(t_1), \dots, j(t_k)) \text{ // by IH on each of } t_1 \dots t_k$$

$$\begin{array}{c} j(\bullet f) \\ / \quad \backslash \\ t_1 \dots t_k \end{array} \text{ // } j \text{ is a } \Sigma\text{-homomorphism from } \text{Tree}_{\Sigma} \text{ to } \mathcal{B}$$

*)

(*)

Variables

We now move from expressions to formulas, i.e., expressions that contain *variables*. As the name implies, a variable is a 0-ary symbol can be given different *values*, in contrast with constants in a signature Σ , whose *values* are fixed, once we fix a Σ -algebra \mathcal{A} . Assume a (denumerable) set \mathcal{X} of variables, disjoint from symbols in Σ , i.e., $\mathcal{X} \cap \Sigma = \emptyset$ *)

(* Let us represent variables using the OCaml type `string` and extend the encoding of the abstract syntax for expressions in the data type `exp`:

*)

```
type exp = N of int
          | V of string
          | Plus of exp * exp
```

```
| Times of exp * exp;;
```

(* Note that we have a new family of base cases: variables, which are denoted using the parameterised constructor `V(_)` which takes arguments of type `string`.

*)

(* The corresponding Concrete Syntax can also be extended

```
E -> T
```

```
E -> T + E
```

```
T -> F
```

```
T -> F * T
```

```
F -> n
```

```
F -> v // a family of variables
```

```
F -> (E)
```

*)

(*

We now extend the syntactic functions on tree-structured expressions *)

(* `size`, as before, returns the number of nodes in the tree, with a variable counted as a single node. *)

```
let rec size t = match t with
  N _ -> 1
| V _ -> 1
| Plus(t1, t2) -> (size t1) + (size t2) + 1
| Times(t1, t2) -> (size t1) + (size t2) + 1
;;
```

(* `ht`, also as before, returns one less than the number of levels in the tree, with variables being treated as leaf nodes. *)

```
let rec ht t = match t with
  N _ -> 0
| V _ -> 0
| Plus(t1, t2) -> (max (ht t1) (ht t2)) + 1
| Times(t1, t2) -> (max (ht t1) (ht t2)) + 1
;;
```

(* As mentioned before, both `size` and `ht` are good measures for performing induction on trees. *)

(* We now define a syntactic function on `exp` called `vars` which returns the *list* of variables in the tree. (We will later see that we really want a function returning the *set* of variables in a syntax tree. If we represent sets as lists, then the difference goes away; of course, we must avoid duplicates and not care about the relative order in which elements appear). *)

```
let rec vars t = match t with
  N _ -> [ ]
| V x -> [x]
| Plus(t1, t2) -> (vars t1) @ (vars t2)
```

```

    | Times(t1, t2) -> (vars t1) @ (vars t2)
;;

```

(* We now extend the *standard semantics* — the definitional interpreter of a language of expressions which include variables. The main question to answer is “what is the **value** of an **expression** that contains variables?” The way to approach this question is to pose a counter-question, asking what **value** each variable takes. So `eval` now takes an extra argument, which we call `rho`, that maps each variable to a value. For the most part, except the new base cases, argument `rho` is carried along as extra baggage in each call. But it cannot be left out. (Why not?) Note that `rho` is a function from `string` to `int`! *)

```

let rec eval t rho = match t with
  | N n -> n
  | V x -> rho x
  | Plus(t1, t2) -> (eval t1 rho) + (eval t2 rho)
  | Times(t1, t2) -> (eval t1 rho) * (eval t2 rho)
;;

```

(* We now extend the `compile` function. First we introduce an opcode `LOOKUP` to look up the value associated with a given variable.

```

*)
type opcode = LD of int
             | LOOKUP of string //look up variable's value
             | PLUS
             | TIMES
;;

```

(* The compiler as before, is a post-order traversal, i.e., a tree-recursive function, with a straightforward new base case. *)

```

let rec compile t = match t with
  | N n -> [LD (n)]
  | V x -> [ LOOKUP(x) ]
  | Plus(t1, t2) -> (compile t1) @ (compile t2) @ [PLUS]
  | Times(t1, t2) -> (compile t1) @ (compile t2) @ [TIMES]
;;

```

(* The stack machine is again written as a tail recursive function driven by the first op-code and the state of the stack, and a new component `env` (environment), which maps variables to their corresponding answers.

```

*)
let rec stkmc s code env = match code with
  | [] -> s
  | (LD(n)) :: code' -> stkmc (n :: s) code' env
  | (LOOKUP(x)) :: code' -> let n=env x
                           in stkmc (n::s) code' env
  | PLUS :: code' -> (match s with
    n1 :: n2 :: s' -> let n = n1 + n2

```

```

                                in stkmc (n :: s') code' env
| _ -> raise Stuck)
| TIMES :: code' -> (match s with
  n1 :: n2 :: s' -> let n = n1 * n2
                    in stkmc (n :: s') code' env
| _ -> raise Stuck)
;;

```

(* Let us now abstract this treatment to arbitrary **abstract syntax trees** over arbitrary **signatures** Σ . We do so by defining a set $Tree_\Sigma(\mathcal{X})$.

Definition: Suppose Σ is a given signature.

Suppose \mathcal{X} is a (denumerable) set of variables.

The set $Tree_\Sigma(\mathcal{X})$ is inductively defined as follows:

Base cases:

- for each 0-ary symbol c in Σ , a labelled node $\bullet c$ is in $Tree_\Sigma(\mathcal{X})$
- for each variable x in \mathcal{X} , a labelled node $\bullet x$ is in $Tree_\Sigma(\mathcal{X})$.

Induction cases: for each $k > 0$,

for each k -ary symbol f in Σ ,

for any trees t_1, \dots, t_k in $Tree_\Sigma(\mathcal{X})$,

the tree consisting of a labelled node $\bullet f$ at the root

with k trees t_1, \dots, t_k ordered as subtrees below the root node

is in $Tree_\Sigma(\mathcal{X})$.

The algebra of abstract syntax trees

$Tree_\Sigma(\mathcal{X})$ is the Σ -algebra with carrier set $Tree_\Sigma(\mathcal{X})$, where (as before)

- every 0-ary symbol c in Σ is interpreted as the labelled node $\bullet c$, in $Tree_\Sigma(\mathcal{X})$,
- every k -ary symbol ($k > 0$) f in Σ is interpreted as the tree-forming function that takes k trees t_1, \dots, t_k from $Tree_\Sigma(\mathcal{X})$, and creates the tree



by sticking t_1, \dots, t_k as the subtrees below a new root node labelled $\bullet f$. *)

(* Observe \mathcal{X} is in 1-1 correspondence with a subset $\bullet \mathcal{X}$ of $Tree_\Sigma(\mathcal{X})$.

Note also that $Tree_\Sigma(\mathcal{X})$ contains $\bullet \mathcal{X} \cup Tree_\Sigma$.

Definition (B-valuation). A function $\rho : \mathcal{X} \rightarrow B$ that maps variables to values in a carrier set B is called a **B-valuation**.

Definition (Function extension)

A function $h : Tree_\Sigma(\mathcal{X}) \rightarrow B$ extends function $\rho : \mathcal{X} \rightarrow B$ if

for all $x \in \mathcal{X}$, $h(\bullet x) = \rho(x)$

Unique Homomorphic Extension Theorem

For any signature Σ , and

any Σ -algebra, $\mathcal{B} = \langle B, \dots \rangle$,

and any B -valuation $\rho : \mathcal{X} \rightarrow B$,

there is a unique Σ -homomorphic extension of $\rho : \mathcal{X} \rightarrow B$, written

$$\hat{\rho} : \text{Tree}_{\Sigma}(\mathcal{X}) \rightarrow B$$

from the (syntactic) Σ -algebra $\text{Tree}_{\Sigma}(\mathcal{X})$ to the Σ -algebra \mathcal{B} .

Proof (First) Existence of a Σ -homomorphic extension of $\rho : \mathcal{X} \rightarrow B$.

By construction: define $\hat{\rho} : \text{Tree}_{\Sigma}(\mathcal{X}) \rightarrow B$ to be a Σ -homomorphism as follows:

- For each 0-ary symbol c in Σ : $\hat{\rho}(\bullet c) = c_{\mathcal{B}}$
- For each $x \in \mathcal{X}$, $\hat{\rho}(\bullet x) = \rho(x)$
- For each k -ary symbol ($k > 0$) f in Σ , and t_1, \dots, t_k in $\text{Tree}_{\Sigma}(\mathcal{X})$

$$\begin{array}{c} \hat{\rho}(\bullet f) = f_{\mathcal{B}}(\hat{\rho}(t_1), \dots, \hat{\rho}(t_k)) \\ \begin{array}{c} / \quad \backslash \\ t_1 \dots t_k \end{array} \end{array}$$

Uniqueness of $\hat{\rho}$

Suppose j is an extension of ρ , and $j : \text{Tree}_{\Sigma}(\mathcal{X}) \rightarrow B$ is a Σ -homomorphism from Σ -algebra $\text{Tree}_{\Sigma}(\mathcal{X})$ to the Σ -algebra \mathcal{B} .

Proof by induction on $(\text{ht } t)$ that for all t in $\text{Tree}_{\Sigma}(\mathcal{X})$, $\hat{\rho}(t) = j(t)$

Base cases $(\text{ht } t) = 0$.

subcase t is of the form $\bullet c$

$$\begin{aligned} \text{for each 0-ary symbol } c \text{ in } \Sigma: \quad \hat{\rho}(\bullet c) &= c_{\mathcal{B}} \quad // \text{ defn of } \hat{\rho} \\ &= j(\bullet c) \quad // j \text{ is a } \Sigma\text{-homomorphism} \\ &\quad // \text{ from } \text{Tree}_{\Sigma}(\mathcal{X}) \text{ to } \mathcal{B} \end{aligned}$$

subcase t is of the form $\bullet x$ for some $x \in \mathcal{X}$

$$\begin{aligned} \hat{\rho}(\bullet x) &= \rho(x) \quad // \text{ defn of } \hat{\rho} \\ &= j(\bullet x) \quad // j \text{ extends } \rho \end{aligned}$$

Induction Hypothesis:

Assume that for all t' in $\text{Tree}_{\Sigma}(\mathcal{X})$ such that $(\text{ht } t') \leq n$,

$$\hat{\rho}(t') = j(t')$$

Induction Step: Consider t in $\text{Tree}_{\Sigma}(\mathcal{X})$ s.t. $(\text{ht } t) = n+1$

t must be of the form

$$\begin{array}{c} \bullet f \\ / \quad \backslash \\ t_1 \dots t_k \end{array}$$

with $(\text{ht } t_i) \leq n$ ($1 \leq i \leq k$)

Now, by definition,

for each k -ary symbol ($k > 0$) f in Σ , and t_1, \dots, t_k in $\text{Tree}_{\Sigma}(\mathcal{X})$ s.t. $(\text{ht } t_i) \leq n$

$$\begin{array}{c} \hat{\rho}(\bullet f) \\ / \quad \backslash \\ t_1 \dots t_k \end{array}$$

$$\begin{aligned}
&= f_{\mathcal{B}}(\widehat{\rho}(t_1), \dots, \widehat{\rho}(t_k)) \quad // \text{definition of } \widehat{\rho} \\
&= f_{\mathcal{B}}(j(t_1), \dots, j(t_k)) \quad // \text{by IH on each of } t_1 \dots t_k \\
&= j(\textstyle \begin{array}{c} \bullet f \\ / \quad \backslash \\ t_1 \dots t_k \end{array}) \quad // \begin{array}{l} j : \text{Tree}_{\Sigma}(\mathcal{X}) \rightarrow B \text{ is a } \Sigma\text{-homomorphism} \\ // \text{from } \Sigma\text{-algebra } \text{Tree}_{\Sigma}(\mathcal{X}) \text{ to the } \Sigma\text{-algebra } \mathcal{B}. \end{array}
\end{aligned}$$

*)

(* Only those variables that appear in a tree (term, expression) matter in its evaluation.

Lemma (Relevant Variables): Suppose Σ is an *arbitrary* signature and \mathcal{X} is the set of variables.

Let t in $\text{Tree}_{\Sigma}(\mathcal{X})$ be an arbitrary tree (term, expression).

Let $V = \text{vars}(t)$.

Consider any Σ -algebra $\mathcal{B} = \langle B, \dots \rangle$, and let $\rho_1, \rho_2 \in \mathcal{X} \rightarrow B$ be any two B -valuations such that for all variables $x \in V$, $\rho_1(x) = \rho_2(x)$.

[Note: $\rho_1(y)$ and $\rho_2(y)$ may be very different for $y \notin V$.]

Then $\widehat{\rho}_1(t) = \widehat{\rho}_2(t)$.

Proof by induction on $\text{ht}(t)$.

Base cases ($\text{ht}(t) = 0$).

subcases: t is of the form $\bullet c$ for some 0-ary symbol c in Σ .

Therefore $\widehat{\rho}_1(\bullet c) = c_{\mathcal{B}} = \widehat{\rho}_2(\bullet c)$

subcases t is of the form $\bullet x$ for some x in \mathcal{X} . So $x \in V$.

Therefore $\widehat{\rho}_1(\bullet x) = \rho_1(x) = \rho_2(x) = \widehat{\rho}_2(\bullet x)$

Induction Hypothesis: Assume for all t' in $\text{Tree}_{\Sigma}(\mathcal{X})$ such that $\text{ht}(t') \leq n$,

for all $\rho'_1, \rho'_2 \in \mathcal{X} \rightarrow B$ such that

for all variables $x \in V' = \text{vars}(t')$, $\rho'_1(x) = \rho'_2(x)$,

that $\widehat{\rho}'_1(t') = \widehat{\rho}'_2(t')$.

Induction Step: Let t in $\text{Tree}_{\Sigma}(\mathcal{X})$ be such that $\text{ht}(t) = n+1$.

By analysis t is of the form $\bullet f$ for some symbol f of arity k in Σ

$$\begin{array}{c} \bullet f \\ / \quad \backslash \\ t_1 \dots t_k \end{array}$$

and trees t_1, \dots, t_k from $\text{Tree}_{\Sigma}(\mathcal{X})$, where $\max \text{ht}(t_i) = n$. ($1 \leq i \leq k$)

$$\widehat{\rho}_1(\bullet f) = f_{\mathcal{B}}(\widehat{\rho}_1(t_1), \dots, \widehat{\rho}_1(t_k))$$

$$\begin{array}{c} / \quad \backslash \\ t_1 \dots t_k \end{array}$$

$$= f_{\mathcal{B}}(\widehat{\rho}_2(t_1), \dots, \widehat{\rho}_2(t_k)) \quad // \text{By IH on } t_1, \dots, t_k \text{ with } \rho'_1 = \rho_1 \text{ and } \rho'_2 = \rho_2$$

// Note that $V_i = \text{vars}(t_i) \subseteq \text{vars}(t) = V$

// so if for all $x \in V$, $\rho_1(x) = \rho_2(x)$,

// then for all $x \in V_i$, $\rho_1(x) = \rho_2(x)$ ($1 \leq i \leq k$)

$$= \widehat{\rho}_2(\bullet f)$$

$$\begin{array}{c} / \quad \backslash \\ t_1 \dots t_k \end{array}$$

*)

(*

Proposition: For any t in $Tree_{\Sigma}(\mathcal{X})$, $V = \text{vars}(t)$ is finite.

Corollary: t in $Tree_{\Sigma}(\mathcal{X})$ can evaluate to at most $|B|^{|V|}$ values.

*)